

[Upgrade](#)[Open in app](#)

Eric Hansen

Dec 21, 2021 · 3 min read · [Listen](#)



# Pickling The Particular Pipelines

## A brief foray into two less-often-inspected SKLearn scenarios: custom features within a Pipeline and how to Pickle that Pipeline.

If you've used and read up on SKlearn, there's a good chance you've seen Pipelines mentioned as a powerful tool for ensuring streamlined workflow, avoiding data leakage, and creating a process flow suitable for serialization using Python's pickle module.

There are some incredible comprehensive guides already out there that give overviews at a variety of levels. Recently, though, I had difficulty finding a guide on how to perform two tasks:

1. My particular model involved creating some custom features (i.e. not a prepackaged transformation) based on multiple existing columns of my dataset. This proved to be non-trivial to fit it into my pipeline.
2. When the pipeline was finished, I wanted to pickle it, as I'd heard was an advantage of enclosing my process-flow in a pipeline. There was a dearth of examples of how to do this.

So, let's approach these two tasks here.

First, begin with a simple dataset, three numeric columns. Suppose the target column is related to the two predictors in a non-linear way.

```
import pandas as pd
import numpy as np

col_d = {'numeric1':[1, 2, 5, 11, 12],
```





```
df_x = df.drop('target', axis=1)
df_y = df['target']

#suppose we wanted to create non-linear columns because we had reason
to believe
# such a relationship held; and perhaps even knew there was the
modulus operation happening
# but suppose we were off by a multiple, let's say 3

def feat_eng(df):
    df['n1sq'] = np.power(df['numeric1'], 2)
    df['n2sqrt'] = np.power(df['numeric2'], 1/2)
    df['mod2'] = 3*((np.power(df['numeric1'], 2) * np.power(
df['numeric2'], 1/2)) % 2)
    return df

df_xe= feat_eng(df_x)
print(df_xe)
```

So, we have added a function to compute some custom-designed features to our dataframe at this point. As shown above, it'd be no problem to simply pass the dataframe through the `feat_eng()` function — in normal workflow. However, if we want to create a pipeline, this creates an obstacle. Each step in a pipeline needs to support `.fit` and `.transform` methods. Fortunately for us, we can write a class wrapper around our `feat_eng()` function (and even inherit from existing classes) to make it work.

```
from sklearn.base import TransformerMixin, BaseEstimator
class customFeats(TransformerMixin, BaseEstimator):
    '''object wrapper for engineered features, suitable for
    pipelining'''
    def transform(self, X):
        X = feat_eng(X)
        return X

    def fit(self, X, y=None):
        return self
```

For our purposes, `.transform` should apply our `feat_eng()` function, and `.fit` should just





Finally(for part 1), we can wrap the entire process in our pipeline. I've included a standard scaler and used a random forest classifier for this example; of course this can be replaced with other transforms and classifiers as necessary.

```
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier

pipe = Pipeline(steps=[
    ('customFeats', customFeats()),
    ('scaler', StandardScaler()),
    ('RFclf', RandomForestClassifier( criterion='gini',
n_estimators=10))
])

#

col_d_test = {'numeric1':[1, 3, 5, 11, 12],
              'numeric2': [144, 36, 225, 121, 81],
              'target':[0, 0, 1, 1, 0 ]}
#note, target = ((n1**2*n2**(1/2))%2)

df_test = pd.DataFrame(col_d_test)
df_x_test = df.drop('target', axis=1)
df_y_test = df['target']

pipe.fit(df_x, df_y)
print('Training set accuracy: {:.4}%'.format(100*pipe.score(df_x,
df_y)))
print('Test set accuracy: {:.4}%'.format(100*pipe.score(df_x_test,
df_y_test)))
```

Great! Onto part 2 — putting the newly created pipeline into python's pickle function. This will likely be straightforward, using the 'with' treatment to ease the file I/O handling.

```
import pickle
with open('clf_model_1.pkl', 'wb') as f:
    pickle.dump(pipe,f)
```



[Upgrade](#)[Open in app](#)

```
if unpickled_pipe.score(df_x_test, df_y_test) == pipe.score(df_x_test,
df_y_test):
    print('Pickled and unpickled models produce same scores!')
```

This works great for me, but some users do run into some difficulty with the pipe step and their own custom feature functions. One solution involves saving the `feat_eng()` function as a .py file, then importing it and saving it using `joblib`, but the details of that are beyond the scope of this blog entry.

In closing, pipeline functionality is tremendously useful and there are many great guides out there for it. I hope that this provided you with some useful examples that were otherwise hard to find!

Following is a github gist with the full code from this post:





Upgrade

Open in app

